

AIRA K. SALISH

Chennai, India | +91 70343 77607 | airasalish@gmail.com | linkedin.com/in/airasalish | github.com/airasalish
Business Technology Solutions Associate | AI Engineering, LLM Applications, Evaluation and Analytics

PROFESSIONAL SUMMARY

B.Tech Computer Science and Engineering (Artificial Intelligence and Machine Learning), 2027 batch, who has **built and shipped LLM-based solutions in production** and the **evaluation that proves they work**. Indian citizen, single degree, **no active backlogs**.
Evaluation mindset, applied: designed **evaluation prompts, rubric-based test sets and labelled regression suites** to check LLM outputs against expected behaviour, defined the quality metric a team measured each release against, and **cut incorrect production responses by 35%**.
Prompting and agent engineering: **few-shot prompting, chain-of-thought and structured outputs**, with **bounded agent loops and tool calling**; treats **hallucination** as a measurable failure mode, and disabled a shipped feature after validation showed an output that could not be trusted.
Software engineering and analytics foundation: **Python with OOP, version control, REST APIs** and readable, testable code; **SQL optimised for 30% better reporting efficiency** with **BI dashboards** built for non-technical commercial teams. **AWS Certified AI Practitioner, 1000/1000**.

EDUCATION

SRM Institute of Science and Technology, Chennai, India Expected 2027
B.Tech, Computer Science and Engineering (Artificial Intelligence and Machine Learning) | **CGPA: 7.8 / 10 (78%)** | Class XII: **72%** | Class X: **90%** | No active backlogs | Single degree, not dual or Master's | Indian citizen
Coursework: Machine Learning, Data Science, Deep Learning fundamentals including transformer architecture, attention and tokenization, Data Structures and Algorithms, Object-Oriented Programming, Database Management Systems, Operating Systems, Software Engineering and Project Management

TECHNICAL SKILLS

Large Language Models and Prompting: LLM primitives — tokens, embeddings, context windows, sampling parameters including temperature and top-p; few-shot prompting, chain-of-thought prompting, structured outputs, system prompt design, prompt iteration
AI Agents and Tool Use: Agent loop design, bounded agent loops, tool calling, LLM API integration, retrieval grounding, rate-limit-aware concurrent batching, idempotent and resumable runs
LLM Evaluation and Safety: Evaluation prompts and test-case design, rubric-based test sets, labelled regression suites, held-out evaluation against a baseline, precision and recall, hallucination identification and mitigation, prompt injection awareness, avoiding sensitive data exposure in prompts and logs
Software Engineering: Python including OOP, Java, JavaScript, version control with Git and GitHub, REST API design and integration, readable and testable code, unit and integration testing, debugging and root-cause analysis, Docker, CI/CD, Linux
Data, BI and Analytics: SQL querying, generation and optimisation, data pipelines, data profiling and validation, BI dashboards and reporting, information management, PostgreSQL, MySQL, MongoDB, pandas, NumPy, Jupyter
Deep Learning and ML: Transformer architecture, attention, tokenization, LLM training pipeline concepts, supervised learning, model evaluation metrics, trade-offs between retrieval and fine-tuning
Delivery and Consulting Skills: Agile and Waterfall delivery, requirement gathering from non-technical stakeholders, solution design and prototyping, test case authoring, technical documentation, proactive risk flagging, stakeholder communication
Currently developing: Kubernetes and container orchestration, C programming

PROFESSIONAL EXPERIENCE

Generative AI Researcher Dec 2025 - Feb 2026 Remote
The Spotted App, AI product team under Amit Baradia, BeSuperhuman.ai

- **Built the evaluation framework for a production LLM system:** wrote **evaluation prompts and rubric-based test sets**, maintained **labelled regression suites** and reproducible failure cases, and **defined the quality metric** every release was measured against — **reducing incorrect production responses by 35%**.
- **Identified and documented recurring failure modes including hallucination**, published them in an **operating playbook adopted team-wide**, and tracked each defect through reproduction, root-cause analysis and **verified closure**.
- **Communicated proactively with founders**, reporting progress, findings and **risks early**, and carried their feedback into the next delivery cycle.

Prompt Engineering and Automation Intern May 2025 - Sept 2025 Bengaluru, Hybrid
BabyBillion, kids content platform, 1.5B+ monthly views; under Dinesh Godra, ex-VP Unacademy

- **Engineered prompts and generative-AI pipelines** for production content, iterating on structured outputs against measurable targets; **one generative-AI short reached 2 crore (20 million) views**.
- **Automated repetitive delivery work** with Python across 50+ concurrent workflows, **removing 40% of the team's recurring manual effort**.
- **Optimised SQL for 30% better reporting efficiency** and built **self-serve BI dashboards** so non-technical commercial teams answered routine questions without raising a request; **profiled and validated data** and escalated quality defects before they reached business reporting.

PROJECTS

WARRANT, LLM Agent for Automated Chargeback Risk Decisions | Python | Razorpay AI Buildathon 2026 Sept 2026
github.com/airasalish/warrant---Razorpay-2026

- Built a **bounded agent loop around LLM inference** with the **risk signals computed deterministically in code**, so automated decisions are **reproducible and auditable** rather than dependent on model variance. Zero false positives and zero false negatives on committed decisions; **100% defence rate under adversarial testing**, including prompt-injection style inputs.
- **Evaluated on a held-out test set rather than development data and published the weaker result** — 71% recall held-out against 92% on development — with limitations documented openly. 51 commits, public repository.

DRIFT, Deployed Data Platform with Rule-Based Signals and Audit Trail | Python, PostgreSQL, React 2026
github.com/airasalish/drift | [live demo](#) | [technical video](#)

- Designed and deployed the **full stack end to end** — ingestion, database, backend, interface — with **five transparent rules computed in code**, each surfacing its evidence, and an **optional AI summary layer kept separate from the deterministic logic**.
- **191 commits, 121 automated tests all passing**, covering stale and failed-fetch paths, plus a dated log of **20+ design decisions and production defects found and fixed**.

HackerRank Orchestrate, Global AI Systems Competition | Python, LLM APIs Jun - Sept 2026

- **Improved evaluated accuracy from 79% to 87%** through baseline measurement, systematic error analysis and re-measurement after each prompt and pipeline change. Peak rank **#27 of 12,885+ participants**; **defended the design in a live 30-minute technical interview**.

CERTIFICATIONS

[AWS Certified AI Practitioner, 1000/1000 \(2026\)](#) | [Red Hat Certified Specialist in Enterprise Linux Automation, RH294 \(2026\)](#) | [SAP Certified Back-End Developer, ABAP Cloud \(2026\)](#)

ACHIEVEMENTS AND LEADERSHIP

Invited to **HackerRank Earn**, an invitation-only expert network (2026) | Top 5% of 1,200+ applicants, **CitiBridge 2026, Citibank** | State Finalist, OpenAI and NxtWave Buildathon | Contributor, GirlScript Summer of Code 2025
Member, **ACM-W SRMIST**, the ACM Women's chapter at SRM, 120+ members; Technical Lead, AI/ML (2024 to 2026)